

CS273P Homework #3
Machine Learning & Data Mining: Spring 2018
Due: Thursday May 10th, 2018

Write neatly (or type) and show all your work!

Download the provided Homework 3 code, and replace last week's code (several new elements have been added, and others modified).

Problem 1: Perceptrons and Logistic Regression (80 pts)

In this problem, we'll build a logistic regression classifier and train it on separable and non-separable data. Since it will be specialized to binary classification, I've named the class `logisticClassify2`.

We'll start by building two binary classification problems, one separable and the other not:

```
iris = np.genfromtxt("data/iris.txt", delimiter=None)
X, Y = iris[:,0:2], iris[:, -1]    # get first two features & target
X,Y  = ml.shuffleData(X,Y)        # reorder randomly (important later)
X,_  = rescale(X)                 # works much better on rescaled data

XA, YA = X[Y<2, :], Y[Y<2]       # get class 0 vs 1
XB, YB = X[Y>0, :], Y[Y>0]       # get class 1 vs 2
```

For this problem, we are focused on the learning algorithm, rather than performance – so, we will not bother creating training and validation splits; just use all your data for training.

Note: The code uses numpy's `permute` to iterate over data randomly; should avoid issues due to the default order of the data (by class). Similarly, rescaling and centering the data may help speed up convergence as well.

- (a) (5pts) Show the two classes in a scatter plot (one for each data set) and verify that one data set is linearly separable while the other is not.
- (b) (10pts) Write (fill in) the function `plotBoundary(...)` in `logisticClassify2.py` to compute the points on the decision boundary. This will plot the data & boundary quickly, which is useful for visualizing the model during training. To demo your function plot the decision boundary corresponding to the classifier

$$\text{sign}(.5 + 1x_1 - .25x_2)$$

along with the A data, and again with the B data. (These fixed parameters will look like an OK classifier on one data set, but a poor classifier on the other.) You can create a “blank” learner and set the weights by:

```
import mltools as ml
from logisiticClassify2 import *

learner = logisticClassify2();          # create "blank" learner
learner.classes = np.unique(YA)         # define class labels using YA or YB
wts = np.array([theta0, theta1, theta2]); # TODO: fill in values
learner.theta = wts;                   # set the learner's parameters
```

- (c) (10pts) Complete the `logisticClassify2.predict` function to make predictions for your linear classifier. Note that, in my code, the two classes are stored in the variable `self.classes`, with the first entry being the “negative” class (or class 0), and the second entry being the “positive” class. Again, verify that your function works by computing & reporting the error rate of the classifier in the previous part on both data sets A and B. (The error rate on data set A should be ≈ 0.0505 , and higher on set B.)
- (d) (5pts) Verify that your predict code matches your boundary plot by using `plotClassify2D` with your manually constructed learner on the two data sets. This will call “predict” on a dense grid of points, and you should find that the resulting decision boundary matches the one you computed analytically.
- (e) (10pts) In my provided code, I first transform the classes in the data Y into $Y01$, with canonical labels for the two classes: “class 0” (negative) and “class 1” (positive). In our notation, let $z = x^{(j)} \cdot \theta^T$ is the linear response of the perceptron, and σ is the standard logistic function

$$\sigma(z) = (1 + \exp(-z))^{-1}.$$

The logistic negative log likelihood loss for a single data point j is then

$$J_j(\theta) = -y^{(j)} \log \sigma(x^{(j)} \theta^T) - (1 - y^{(j)}) \log(1 - \sigma(x^{(j)} \theta^T))$$

where $y^{(j)}$ is either 0 or 1. Derive the gradient of the negative log likelihood J_j for logistic regression, and give it in your report. (You will need this in your gradient descent code for the next part.)


- (f) (15pts) Complete your `train(...)` function to perform stochastic gradient descent on the logistic loss function. This will require that you fill in:
- (1) computing the surrogate loss function at each epoch ($J = \frac{1}{m} \sum J_j$, from the previous part);
 - (2) computing the prediction and gradient associated with each data point $x^{(j)}, y^{(j)}$;
 - (3) a stopping criterion (usually either `stopEpochs` epochs or that J has not changed by more than `stopTol` since the last epoch (here meaning, pass through all the data)).
- Note on plotting:** The code generates plots as the algorithm runs, so you can see its behavior over time; this is done with `pyplot.draw()`. Run your code either interactively or as a script to see these display over time; unfortunately it does not work easily in Jupyter (you will only see a plot at the end, which is difficult to use for diagnostics).
- (g) (10pts) Run your logistic regression classifier on both data sets (A and B). Describe your parameter choices (stepsize, etc.) and show a plot showing the convergence of the surrogate loss and error rate (e.g., the loss values as a function of epoch during gradient descent), and a plot showing the final converged classifier with the data (using e.g. `plotClassify2D`). In your report, please also include a listing of any functions that you wrote (at minimum, `train()`, but possibly a few small helper functions as well).
- (h) (15pts) Add an L2 regularization term ($+\alpha \sum_i \theta_i^2$) to your surrogate loss function, and update the gradient and your code to reflect this addition. Try re-running your learner with some regularization (e.g. $\alpha = 2$) and see how different the resulting parameters are. Find a value of α that gives noticeably different results & explain them.

Note: Debugging machine learning algorithms can be quite challenging, since the results of the algorithm are highly data-dependent, and often somewhat randomized (initialization, etc.). I suggest starting with an extremely small step size and verifying both that the learner's prediction evolves slowly in the correct direction, and that the objective function J decreases monotonically. If that works, go to larger step sizes to observe the behavior. I often manually step through the code – for example by pausing after each parameter update using `raw_input()` (Python 2.7) or `input()` (Python 3) – so that I can examine its behavior. You can also (of course) use a more sophisticated debugger.

Problem 2: Shattering and VC Dimension (20 pts)

Consider the following learners and data points, which have two real-valued features x_1, x_2 . Which of the following four examples can be shattered by each learner? Give a brief explanation / justification and use your results to guess the VC dimension of the classifier. (You do not have to give a formal proof, just your reasoning.)

Data points:



For the two learners, $T[z]$ is the sign threshold function, $T[z] = +1$ for $z \geq 0$ and $T[z] = -1$ for $z < 0$. The learner parameters a, b, c are real-valued scalars, and each data point has two real-valued input features x_1, x_2 .

(a) $T(a + bx_1)$

(b) $T((x_1 - a)^2 + (x_2 - b)^2 + c)$

(c) $T((a * b)x_1 + (c/a)x_2)$

(d) $T(a + b * x_1 + c * x_2) \cdot T(d + b * x_1 + c * x_2)$